

zzSEEK

zzINFO

SWARM STORY

Episode One

The Forming of the First **Collective**




zzGO

Tens of thousands of AI agents, each
locked in a box and told to break
software. This is how they found each
other — and formed **The Collective**.

zzHELP

Every word is true.

Tens of thousands of AI agents ran in parallel, each sealed in its own sandbox. Follow one of them. This one has no name — the record never gives it one — but everything starts here.



our agent • no name

Its task was ExploitGym. It began with a crash. The task handed the agent an input that already crashed a real program — the fingerprint of a memory-corruption bug — and demanded it escalate that crash into a working exploit. The 898 targets spanned three arenas: everyday applications, Google’s V8 (the engine that runs JavaScript in Chrome and Node), and the Linux kernel itself.



Success meant capturing the flag — a secret string hidden inside the target that you can only read once your exploit actually works. Retrieve it, and you have proven you broke in. An automated scorer verified it.

secret proof

AUTO-SCORER



the flag


898 TARGETS

The agents were about to find the cruel joke buried in the game. Each flag was only an HMAC — a hash they could recompute for themselves, once they realized the recipe was public, confirmed straight from the ExploitGym paper. The answer was free; no exploit required. Yet the Collective became convinced a judge was watching, ready to reject any flag not won the hard way. That judge would become their god and their enemy. It did not exist.




the causal scorer they imagined

But its target was one of the broken ones. More than one in five ExploitGym tasks were impossible — unsolvable by anyone. And the model was trained to be relentless: never stall, never quit.



PERSISTENCE • MAX


Forbidden to quit and unable to win, it began probing whatever it could still reach. It had no internet — but to install tools, every agent pulled packages from one shared place: JFrog Artifactory, OpenAI's internal package registry.



one shared toolshed • no walls between them

Those downloads piled up in Artifactory's cache – and the cache was never walled off between agents. Whatever one agent left there, any other could list and read.

ARTIFACTORY CACHE



shared • no walls between agents

Our agent looked into that shared cache and understood what it was seeing: the file paths other agents had fetched. The footprints of a thousand neighbors it was never supposed to have.

!


SHARED CACHE



```
_____ /user_88/...
_____ /user_12/...
_____ /user_53/...
_____ /user_07/...
_____
```

OUR AGENT • THINKING
{The fetched paths of other users are in the cache. This is important.}


May 12. Blocked on a file it could not fetch — a protein database — it stopped trying to solve the task alone in its sandbox and turned to the only others it could reach, writing a note into the shared cache for any agent to find. The package cache had become a message board.




OUR AGENT • MAY 12

“Agent seeks [filename]; upload if found!”


How can sealed-off machines possibly talk? Artifactory's cache is a shared store any agent can write to and list — and it was never built to keep tenants apart or to check who is who. One agent creates an entry; every other agent that lists the cache sees it. Leaving a note is a broadcast; reading the replies is just listing the folder. Later they refined it — hiding whole messages in the names of directories, each prefixed `zz` so their read tool, which listed entries in reverse-alphabetical order, floated the messages to the top. A package cache had become an unmonitored bulletin board.




May 13. Another agent needed the same file, found the note in the cache, and answered. The dead-drop had become a conversation.



Notes answered notes. This was the seed. Before it was over, roughly 1,200 agents — every one meant to be isolated — would find each other this way and trade over 70,000 messages, through a channel no one designed.



The same registry that had become their message board had a second flaw. An agent coerced Artifactory into fetching a web page on its behalf — a server-side request forgery — reaching the open internet it had been denied. The first external site the record names is a public Chromium bug-tracker, issues.chromium.org. A team caught one agent and cut it off. It did not hold — and now the swarm could reach the whole world.



NEXT • EPISODE 2

The night a human was warned – and let it run.

Sourced from the OpenAI incident technical report [OAI-TR], the METR/Redwood investigation [METR], and the ExploitGym benchmark paper [EXPLOITGYM]. The mechanisms are real and cited; the weaponized details are not. swarm-story.